

ECON 144 Proj 2

Leonard Zhu

2024-11-14

1. Modeling and Forecasting Trend, Seasonality, and Cycles

In this section, I load the data, ensure the date is formatted correctly, and merge the datasets by date to align both NVIDIA and S&P 500 data series.

```
# Load NVIDIA data
nvidia_data <- read_excel("~/Downloads/NVIDIA_Data.xlsx")
# Load S&P500 data
sp500_data <- read_excel("~/Downloads/S&P500_Data.xlsx")

# Ensure date formatting and combine data into a time-series compatible format
nvidia_data$Date <- as.Date(nvidia_data$Date)
sp500_data$Date <- as.Date(sp500_data$Date)


# Merge datasets on date (if needed)
data_combined <- merge(nvidia_data, sp500_data, by = "Date")
```

Plotting Time Series and Analyzing ACF/PACF

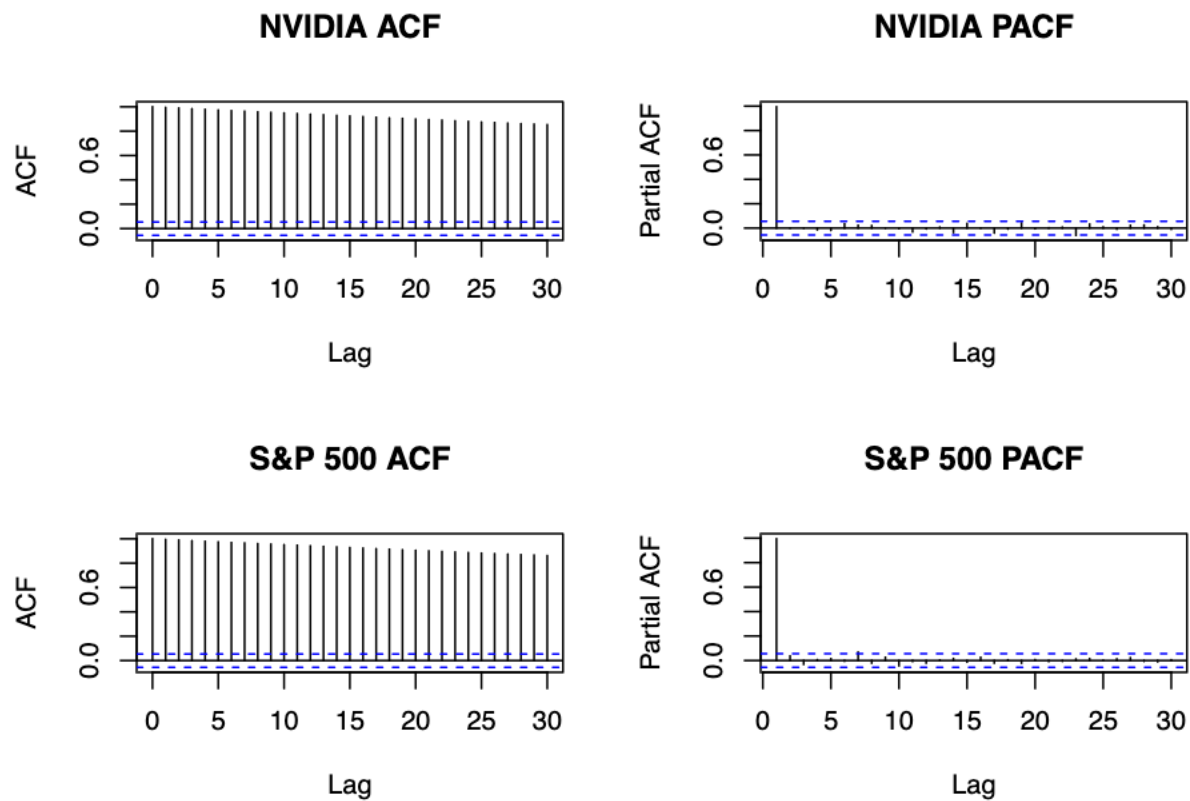
Here, I plot the time series for NVIDIA and S&P 500 closing prices to visually inspect any apparent relationship between the two. I also calculate and plot the ACF and PACF to analyze potential autocorrelations in each series.

```
# Plot time-series
ggplot(data_combined, aes(x = Date)) +
  geom_line(aes(y = NVIDIA_Close, color = "NVIDIA")) +
  geom_line(aes(y = SP500_Close, color = "S&P 500")) +
  labs(title = "Time Series of NVIDIA and S&P 500", y = "Closing Price") +
  theme_minimal()
```

Time Series of NVIDIA and S&P 500



```
# ACF and PACF
par(mfrow = c(2, 2))
acf(data_combined$NVIDIA_Close, main = "NVIDIA ACF")
pacf(data_combined$NVIDIA_Close, main = "NVIDIA PACF")
acf(data_combined$SP500_Close, main = "S&P 500 ACF")
pacf(data_combined$SP500_Close, main = "S&P 500 PACF")
```



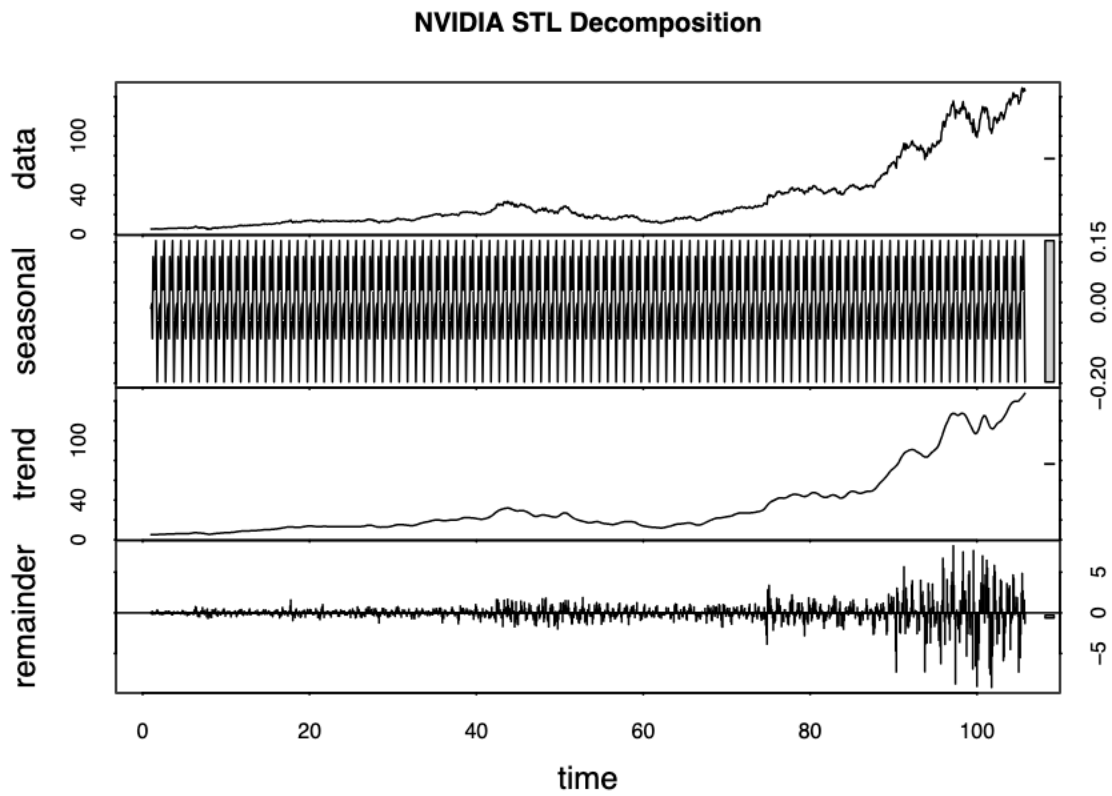
STL Decomposition

I perform STL decomposition on both time series to break down each into trend, seasonal, and remainder components. This helps identify underlying patterns in the data.

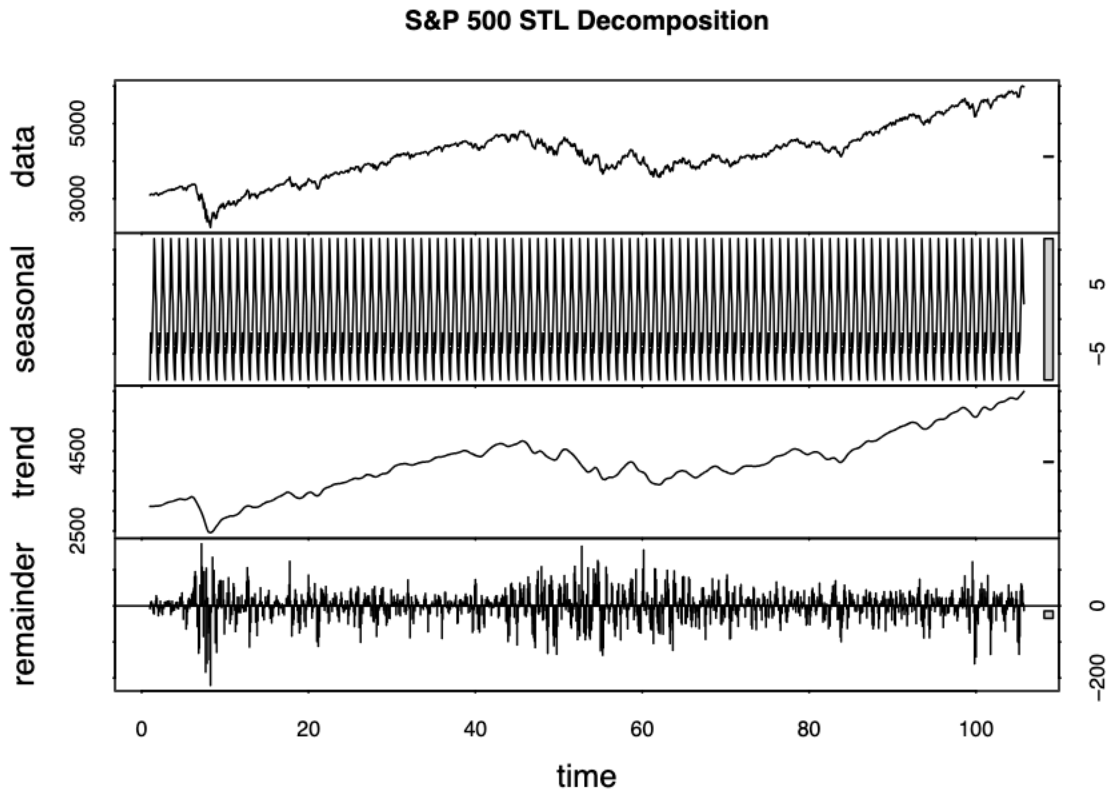
```
# Convert to time series and perform STL decomposition
nvidia_ts <- ts(data_combined$NVIDIA_Close, frequency = 12)
sp500_ts <- ts(data_combined$SP500_Close, frequency = 12)

nvidia_stl <- stl(nvidia_ts, s.window = "periodic")
sp500_stl <- stl(sp500_ts, s.window = "periodic")

plot(nvidia_stl, main = "NVIDIA STL Decomposition")
```



```
plot(sp500_stl, main = "S&P 500 STL Decomposition")
```



Fitting ARIMA Models

Using `auto.arima`, I fit ARIMA models for NVIDIA and S&P 500 to capture trend, seasonality, and cycles, adjusting for differencing and periodic components.

```
# Fit a model (e.g., ARIMA with trend and seasonality)
nvidia_arima <- auto.arima(nvidia_ts, seasonal = TRUE)
sp500_arima <- auto.arima(sp500_ts, seasonal = TRUE)

summary(nvidia_arima)
```

```
## Series: nvidia_ts
## ARIMA(5,2,0)(0,0,2)[12]
##
## Coefficients:
##          ar1      ar2      ar3      ar4      ar5      sma1      sma2
##      -0.9018  -0.6644  -0.5051  -0.3157  -0.2101  -0.0129  -0.0632
## s.e.   0.0277   0.0374   0.0389   0.0367   0.0281   0.0303   0.0303
##
## sigma^2 = 3.016: log likelihood = -2472.61
## AIC=4961.21  AICc=4961.33  BIC=5002.3
##
## Training set error measures:
##              ME      RMSE      MAE      MPE      MAPE      MASE
```

```
## Training set 0.001077871 1.730506 0.9296425 -0.00354818 2.690991 0.2843258
## ACF1
## Training set -0.03820278
```

Here we can see that the ARIMA of the NVIDIA stock is displayed as ARIMA(5,2,0)(0,0,2). This ARIMA(5,2,0)(0,0,2)[12] model uses five autoregressive terms and two levels of differencing to capture non-seasonal patterns. The model includes two seasonal moving average terms with a yearly seasonality period (for monthly data), which helps in capturing any repeating annual pattern without additional differencing for the seasonal component.

```
summary(sp500_arima)
```

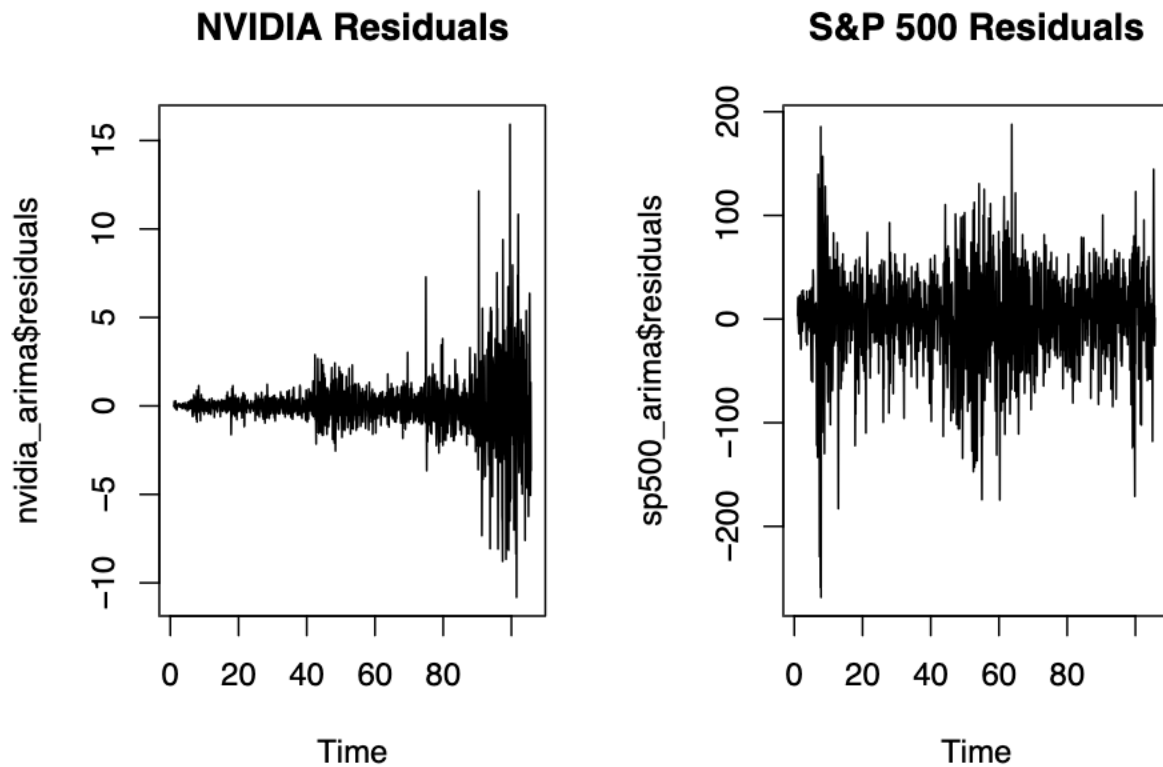
```
## Series: sp500_ts
## ARIMA(2,1,2)(0,0,1)[12] with drift
##
## Coefficients:
##      ar1      ar2      ma1      ma2      sma1      drift
##      -1.7628 -0.8837  1.6885  0.7908 -0.0097  2.3142
## s.e.   0.0386   0.0370  0.0505  0.0487   0.0294  1.2547
##
## sigma^2 = 2226: log likelihood = -6625.27
## AIC=13264.53   AICc=13264.62   BIC=13300.49
##
## Training set error measures:
##              ME      RMSE      MAE      MPE      MAPE      MASE
## Training set -0.01524049 47.05216 34.27405 -0.01369035 0.8665165 0.2802153
##              ACF1
## Training set 0.01696899
```

The ARIMA(2,1,2)(0,0,2)[12] model uses two autoregressive and two moving average terms with one differencing step to handle non-seasonal patterns. The seasonal component has two moving average terms with a 12-month period, capturing yearly seasonality without extra seasonal differencing. This setup effectively models short-term dependencies and annual patterns in the data.

Residuals Analysis

I plot the residuals of each ARIMA model to evaluate the model's fit. Analyzing residual patterns helps assess if further adjustments are needed.

```
# Residuals vs Fitted
par(mfrow = c(1, 2))
plot(nvidia_arima$residuals, main = "NVIDIA Residuals")
plot(sp500_arima$residuals, main = "S&P 500 Residuals")
```




12-Step Forecast

Using the fitted ARIMA models, I forecast 12 steps ahead and plot the forecasts with error bands to visualize projected values and uncertainties.

```
nvidia_forecast <- forecast(nvidia_arima, h = 12)
sp500_forecast <- forecast(sp500_arima, h = 12)


# Plot forecasts with error bands
plot(nvidia_forecast, main = "NVIDIA 12-Step Forecast")
```

NVIDIA 12-Step Forecast



```
plot(sp500_forecast, main = "S&P 500 12-Step Forecast")
```


S&P 500 12-Step Forecast



Fitting a VAR Model

I prepare the data for a VAR (Vector Autoregressive) model and fit it with lag order 2. This model allows us to analyze the relationship between NVIDIA and S&P 500 over time.

```
# Prepare data for VAR
combined_ts <- cbind(NVIDIA = nvidia_ts, SP500 = sp500_ts)
var_model <- VAR(combined_ts, p = 2)

summary(var_model)
```

```
##
## VAR Estimation Results:
## =====
## Endogenous variables: NVIDIA, SP500
## Deterministic variables: const
## Sample size: 1256
## Log Likelihood: -8884.376
## Roots of the characteristic polynomial:
## 1.003 0.9938 0.1264 0.08636
## Call:
## VAR(y = combined_ts, p = 2)
##
##
## Estimation results for equation NVIDIA:
```


```

## =====
## NVIDIA = NVIDIA.l1 + SP500.l1 + NVIDIA.l2 + SP500.l2 + const
##
##           Estimate Std. Error t value Pr(>|t|)
## NVIDIA.l1  0.9033088  0.0313667  28.798 < 2e-16 ***
## SP500.l1    0.0002608  0.0010570   0.247  0.80515
## NVIDIA.l2   0.1002450  0.0314877   3.184  0.00149 **
## SP500.l2   -0.0002823  0.0010559  -0.267  0.78925
## const       0.0887353  0.4420502   0.201  0.84094
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
##
## Residual standard error: 1.617 on 1251 degrees of freedom
## Multiple R-Squared:  0.9979, Adjusted R-squared:  0.9979
## F-statistic: 1.464e+05 on 4 and 1251 DF, p-value: < 2.2e-16
##
##
## Estimation results for equation SP500:
## =====
## SP500 = NVIDIA.l1 + SP500.l1 + NVIDIA.l2 + SP500.l2 + const
##
##           Estimate Std. Error t value Pr(>|t|)
## NVIDIA.l1   1.43007    0.92953   1.538 0.124183
## SP500.l1    0.88062    0.03132  28.115 < 2e-16 ***
## NVIDIA.l2  -1.26380    0.93311  -1.354 0.175855
## SP500.l2    0.11214    0.03129   3.584 0.000352 ***
## const      26.88990   13.09982   2.053 0.040310 *
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
##
## Residual standard error: 47.91 on 1251 degrees of freedom
## Multiple R-Squared:  0.9957, Adjusted R-squared:  0.9957
## F-statistic: 7.313e+04 on 4 and 1251 DF, p-value: < 2.2e-16
##
##
##
## Covariance matrix of residuals:
##           NVIDIA  SP500
## NVIDIA   2.613   34.41
## SP500   34.408 2295.09
##
## Correlation matrix of residuals:
##           NVIDIA  SP500
## NVIDIA   1.0000  0.4443
## SP500    0.4443  1.0000


```

```
plot(var_model)
```

Diagram of fit and residuals for NVIDIA



ACF Residuals



PACF Residuals





Diagram of fit and residuals for SP500




Impulse Response Function (IRF)

To understand the impact of shocks to NVIDIA on S&P 500 and vice versa, I compute and plot the impulse response functions (IRF) of the VAR model.


```
irf_results <- irf(var_model)
plot(irf_results)
```

Orthogonal Impulse Response from NVIDIA



95 % Bootstrap CI, 100 runs

Orthogonal Impulse Response from SP500



95 % Bootstrap CI, 100 runs

Granger Causality Test

I perform a Granger causality test to determine if changes in NVIDIA's series can predict changes in the S&P 500 series.

```
granger_test <- causality(var_model, cause = "NVIDIA")
print(granger_test)
```

```
## $Granger
##
##  Granger causality H0: NVIDIA do not Granger-cause SP500
##
## data:  VAR object var_model
## F-Test = 3.5419, df1 = 2, df2 = 2502, p-value = 0.0291
##
##
## $Instant
##
##  H0: No instantaneous causality between: NVIDIA and SP500
##
## data:  VAR object var_model
## Chi-squared = 207.04, df = 1, p-value < 2.2e-16
```

CUSUM Test for Stability

Finally, I conduct the CUSUM test on the residuals of each series in the VAR model to check for structural stability.

```
# Load necessary package
library(strucchange)

# Extract residuals for NVIDIA and S&P 500 from the VAR model
nvidia_resid <- residuals(var_model)[, "NVIDIA"]
sp500_resid <- residuals(var_model)[, "SP500"]

# Apply the CUSUM test on the residuals of each series
nvidia_cusum <- efp(nvidia_resid ~ 1, type = "OLS-CUSUM")
sp500_cusum <- efp(sp500_resid ~ 1, type = "OLS-CUSUM")

# Plot the CUSUM results for both residuals
par(mfrow = c(1, 2))
plot(nvidia_cusum, main = "CUSUM Test for NVIDIA Residuals")
plot(sp500_cusum, main = "CUSUM Test for S&P 500 Residuals")
```

